COM SCI 239

Quantum Programming: Algorithms in PyQuil

May 12, 2020

Author

Arjun Raghavan and Bryan Pan

1   Overview

2   Design

2.1   Mathematical explanation

All four algorithms take as input a function f in the form:

$$f : \{0,1\}^n \to \{0,1\}^m$$

In the case of Deutsch-Josza, Bernstein-Vazirani, and Grover, we have $m = 1$. Simon, on the other hand, has $m = n$.

For Deutsch-Josza, Bernstein-Vazirani, and Simon, the quantum oracle of f, denoted U_f, is defined as:

$$U_f |x\rangle |b\rangle = |x\rangle |b \oplus f(x)\rangle$$

Here, we have $b \in \{0,1\}^m$. The $\oplus$ operator represents bitwise XOR, or equivalently bitwise addition mod 2.

Grover's algorithm makes use of a gate denoted by Z_f, which is defined as:

$$Z_f |x\rangle = (-1)^{f(x)} |x\rangle$$

Notice that this is precisely the application of the phase kickback trick for gates of a form equivalent to U_f where $|b\rangle = |-\rangle$:

$$Z_f |x\rangle |-\rangle = (-1)^{f(x)} |x\rangle |-\rangle$$

Clearly, we can also define Z_f as a quantum oracle for Grover's algorithm in the same way U_f was defined for the three other algorithms. Given that all four algorithms' quantum oracles have the same general form, we created a single function which could generate a quantum oracle and reused it for each program.

Lemma 1. $\sum_{q \in \{0,1\}^k} |q\rangle \langle q| = I_{2^k}$ *where I_{2^k} is the identity matrix in $\mathcal{H}_{2^k}$ (corresponding to k qubits).*

Proof. Let $q \in \{0,1\}^k$. We have $|q\rangle = \begin{bmatrix} q_1 & q_2 & \dots & q_{2^k} \end{bmatrix}^T$ where $q_i = 1$ for some $1 \le i \le 2^k$ and $q_j = 0$ for all $j \ne i$. Then, $|q\rangle \langle q|$ is a $2^k \times 2^k$ matrix of the form $\begin{bmatrix} q_{ab} \end{bmatrix}$ where:

$$q_{ab} = \begin{cases} 1 & \text{if } a = b = i \\ 0 & \text{otherwise} \end{cases}$$

For $\alpha, \beta \in \{0,1\}^k$ such that $\alpha \ne \beta$, we have $|\alpha\rangle \ne |\beta\rangle \implies |\alpha\rangle \langle \alpha| \ne |\beta\rangle \langle \beta|$. Thus:

$$\sum_{q \in \{0,1\}^k} |q\rangle \langle q| = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 \end{bmatrix} + \begin{bmatrix} 0 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 \end{bmatrix} + \dots + \begin{bmatrix} 0 & 0 & \dots & 0 \\ 0 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 \end{bmatrix} = I_{2^k}$$

$\square$

Consider the application of U_f to the outer product of $|x\rangle |b\rangle$ with itself:

$$U_f |x\rangle |b\rangle \langle x| \langle b| = |x\rangle |b \oplus f(x)\rangle \langle x| \langle b|$$
$$= |x\rangle \langle x| |b \oplus f(x)\rangle \langle b|$$

Let $k = n + m$. By taking the sum of this over all bit strings $xb \in \{0,1\}^k$, we get:

$$\sum_{xb \in \{0,1\}^k} U_f \left(|x\rangle |b\rangle \langle x| \langle b|\right) = U_f \left(\sum_{xb \in \{0,1\}^k} |x\rangle |b\rangle \langle x| \langle b| \right)$$
$$= U_f \circ I_{2^k}$$
$$= U_f$$

And so:
$$U_f = \sum_{xb \in \{0,1\}^k} |x\rangle \langle x| |b \oplus f(x)\rangle \langle b| \tag{1}$$

2.2 Python implementation

As shown in Code Block 1, Equation 1 is straightforwardly implemented in Python thanks to the `numpy` library using methods such as `np.kron` (the Kronecker product is a special case of the tensor product over complex matrices and is the nomenclature used by `numpy`) and `np.outer` [1]. Notice in line 155 that `f` is a dictionary, not a function. Indeed, we chose to treat the input function f for all four programs as a dictionary (or a "mapping" as we mostly referred to it) for a few reasons.

For one, this would allow a user more flexibility in providing f as an input to our programs without having to possibly construct tedious `if-elif` statements. In addition, this allowed us to more easily test our programs—we could randomly generate dictionaries representing functions which followed whatever assumptions were required for each algorithm. An example of this for Deutsch-Josza can be seen in Code Block 2.

Our implementation of generating U_f with this generalized method had the additional benefit of making it very easy to parametrize our solutions in n. The `oracle.gen_matrix` method, along with any other methods which require the dimension of the domain and/or range of f, accepts `n` and `m` (as defined previously) as arguments. While `m` is algorithm-dependent and so assigned inside the code, the user can pass in a value for n using the `-num` option[1].

2.2.1 Readability

There were a couple of possible approaches we could have taken to iterating over all $xb \in \{0,1\}^{n+m}$, including nested `for` loops or generating the entire set of bit strings and iterating

---

[1]Discussed in further detail in the README.

```
148        # Accumulator to hold the value of the summation
149        U_f=np.zeros((2**(n+m),2**(n+m)))
150        for xb in range(0, int(2**(n+m))):
151            # Convert index to binary and split it down the middle
152            x  = f'{xb:0{n+m}b}'[:int(n)]
153            b  = f'{xb:0{n+m}b}'[int(n):]
154            # Apply f to x
155            fx = f[x]
156            # Calculate b + f(x)
157            bfx = f'{int(b, 2) ^ int(fx, 2):0{n}b}'
158
159            # Vector representations of x, b, and b+f(x)
160            xv = np.zeros((2**n,1))
161            xv[int(x, 2)] = 1.
162            bv = np.zeros((2**m,1))
163            bv[int(b, 2)] = 1.
164            bfxv = np.zeros((2**m,1))
165            bfxv[int(bfx, 2)] = 1.
166
167            # Accumulate (|x><x| (*) |b + f(x)><b|) into the sum
168            # (*) is the tensor product
169            U_f = np.add(np.kron(np.outer(xv, xv), np.outer(bfxv, bv)), U_f)
```

Code Block 1: Excerpt from the `gen_matrix` function in `oracle.py` showing the implementation of Equation 1 using `numpy`

over that. We eventually settled on the approach seen in Code Block 1, which we felt was more elegant.

The dimension of $\{0, 1\}^{n+m} =: S$ is 2^{n+m}. Furthermore, each element of S is a bitstring which has a decimal equivalent. Thus, it would be just as effective to iterate through each of these decimal equivalents, convert them to binary, and extract x and b, which is precisely what we did. We felt that this was more concise than generating the entirety of S or nesting `for` loops while still maintaining some clarity of our approach, assisting in readability.

In more general terms, we made sure to thoroughly comment our code wherever we felt it was necessary. Our primary functions all have PyDocs (loosely conforming to PEP/8 guidelines [2]), and further comments are added when specific lines of code require further clarification. Overall, though, our code is self-documenting as much as possible. We have also taken care to maintain limited line lengths (keeping 80 characters as a soft limit) and proper indentation (although the Python language just about forces the latter).

```
43      if algo is Algos.DJ:
44          # Constant: f(x) returns 0 or 1 for all x
45          if func is DJ.CONSTANT:
46              val = np.random.choice(['0','1'])
47              oracle_map = {i: val for i in qubits}
48          # Balanced: f(x) returns 0 or 1 for all x
49          # val1 represents set of x that f(x) = 1
50          # val0 represents set of x that f(x) = 0
51          elif func is DJ.BALANCED:
52              val1 = random.sample(qubits, k=int(len(qubits)/2))
53              val0 = set(qubits) - set(val1)
54              oracle_map = {i: '1' for i in val1}
55              temp = {i: '0' for i in val0}
56              oracle_map.update(temp)
```

Code Block 2: Excerpt from the `init_bit_mapping` function in `oracle.py` showing how we randomly generate a function for testing Deutsch-Josza.

3   Evaluation

3.1   Shared code

As opposed to copying identical sections of code between each of our four programs, we opted to have a central module, `oracle.py`, which contains common code relating to the generation of functions f (bit mappings, to be precise) and quantum oracles, which was discussed in the previous section.

Of course, there is still some duplication of code between programs. Each program has a `getUf` function (named `getZf` for Grover) which checks if a U_f matrix for the desired value of n has already been generated and stored[2], in which case it loads it from its file; otherwise, it simply generates a new U_f. All the implementations of this function are identical.

Also, our programs use the `argparse` library [3] to read and process user-defined options[3]. The code for doing so is largely the same across programs, save for some minor algorithm-specific differences (like the names).

Overall, we estimate that each program file has an average of approximately 38 lines of code that is shared (within a reasonable threshold of a few characters' difference) with the other files. Excluding `oracle.py`, the average length of a single program's Python file is 122 lines. As such, the percentage of shared code between programs is approximately 31%.

There are certainly avenues for reducing code reuse. We began to adopt a `class` structure for the entire project, having each algorithm inherit from an `Algo` class which itself had methods common to all four programs. This can in fact be seen in the code for `dj.py` and `bv.py`.

---

[2]This generation and storing of matrices ahead of actually running the algorithms is discussed in the README.

[3]Also discussed in the README.

However, this was abandoned mostly due to the fact that the assignment required individual files for each program, and so this would end up unnecessarily increasing the amount of code reused between files.

3.2   Testing

As mentioned earlier, our implementation of generating and storing quantum oracle gate matrices and usage of `argparse` to handle user input greatly facilitated testing. We quickly realized that compiling and running our code was (naturally) CPU-intensive, so we also used a couple of Oracle Cloud Infrastructure Compute instances so that we could run our code while freeing up our own computers for further work.

For our implementation of Grover, we configured the program such that it would automatically calculate the minimum number of trials required to minimize error based on the size of the input using the equation $k = \lfloor \pi\sqrt{N}/4 \rfloor$ with $N = 2^n$. For our implementation of Simon, given that a full "iteration" was $n - 1$ applications of the quantum oracle (to generate $n - 1$ values for y), we multiplied the user's input number of trials by 4 to minimize error using the equation $\mathcal{P}(\text{not linearly independent}) = \exp(-t/4)$ with t being the number of trials.

We tested our code with a variety of values of n (the length of the input bit string). For Simon and Grover, we also tested with differing values of t (the number of trials); this was unnecessary for Deutsch-Josza and Bernstein-Vazirani, which are deterministic algorithms. We found that our results very closely matched the expected outcomes of applying our algorithms.

For each value of n up to about 7-9 depending on the algorithm, we generated randomized functions (following the appropriate constraints for its corresponding algorithm) and subsequently U_f matrices, storing them locally for later use. We limited n to 7-9 mostly because any values greater than those resulted in massive U_f matrices (the file size of a 10-qubit U_f matrix stored as an `.npy` array exceeded 100 MB). This proved to be a reasonable restriction, because none of our programs were able to operate on more than 7 qubits before the PyQuil compiler crashed or the program itself exceeded its allocated heap space.

For smaller n, however, we were able to make some observations:

- Deutsch-Josza was the fastest algorithm of the four, able to reach 7 qubits before crashing.

- Whether the input function to Deutsch-Josza was balanced or constant was a large factor in its compilation and execution time. For constant functions, the quantum oracle was simply a large identity matrix, and it appears that PyQuil optimizes for such a case, as can be seen below.

- Despite being probabilistic, Simon's and Grover's algorithms are in fact quite accurate even for a smaller number of trials. Of course, this is in an ideal setting on a Quantum Virtual Machine without noise.

- Compilation time is orders of magnitude larger than execution time. All of the programs' execution time were under 10 seconds; compilation time, on the other hand, increased dramatically as n increased—for Simon's algorithm, an input of $n = 4$ had a compilation time of a whopping 36 minutes. It would have been beneficial to compile ahead of time,

but this was not feasible for two reasons: a new gate had to be used for each value of n, meaning the quantum circuit would have to be recompiled for each input regardless, and it was not clear whether we could compile a quantum program and store it as a file.

- Simon's algorithm is the slowest to compile, crashing `quilc` at only 5 input qubits. This is expected behavior; while the other three algorithms have one additional helper qubit, Simon's algorithm must use n additional qubits. Thus, increasing the input bitstring by 1 causes the resulting quantum oracle matrix to balloon to $2^4 = 16$ times its original size.

3.3   Scalability

While as mentioned above we used Oracle Cloud Infrastructure, we were limited to 2 CPU cores on each of our instances. So, one of our local machines was used to perform tests for execution time. Said machine has an `8-core Intel(R) Core(TM) i7-8550U @ 1.80GHz` and a RAM of `16GiB System Memory` (x2 `8GiB SODIMM DDR4 Synchronous Unbuffered 2400 MHz`).

We measured both the compilation as well as execution time of each run for the four algorithms. The execution time was obtained from the logged output of `qvm -S`. The compilation time was measured simply by measuring the system clock before and after calling `QuantumComputer.run_and_measure`[4].

4   Reflection on PyQuil

Our brief introduction to PyQuil showed us that there were a handful of gates which were provided as part of the library. A such, we expected that it would be challenging to create our own gates. Fortunately, we discovered that it was in fact very simple. We had to generate our own matrix, which was straightforward with the help of `numpy`, but the creation of an actual quantum gate was quite literally only two lines (see Code Block 3).

Some pros of the PyQuil documentation include:

- It provides a fine introduction to PyQuil and quantum programming.

Some cons of the documentation are:

- TODO

TODO: Answer the following:

- (IP) List three aspects of quantum programming in PyQuil that turned out to be easy to learn and list three aspects that were difficult to learn.

- List three aspects of quantum programming in PyQuil that the language supports well and list three aspects that PyQuil supports poorly.

- Which feature would you like PyQuil to support to make the quantum programming easier?

---

[4]A more accurate measure of this would be to subtract execution time from the total measured time to "run and measure", but as mentioned before, compilation time vastly dwarfs execution time.

Figure 1: label

- (IP) List three positives and three negatives of the documentation of PyQuil.

- In some cases, PyQuil has its own names for key concepts in quantum programming. Give a dictionary that maps key concepts in quantum programming to the names used in PyQuil.

- How much type checking does the PyQuil implementation do at run time when a program passes data from the classical side to the quantum side and back?

Figure 2: label

```
57    U_f_def = DefGate('U_f', getUf(n, reload))
58    U_f = U_f_def.get_constructor()
```

Code Block 3: The two lines needed to create a quantum gate that can be used in a PyQuil program.

References

[1] S. van der Walt, S. C. Colbert, and G. Varoquaux. The NumPy array: A structure for efficient numerical computation. Computing in Science & Engineering, 13(2):22–30, 2011.

[2] Guido van Rossum, Barry Warsaw, and Nick Coghlan. Style guide for Python code. PEP 8, 2001.

[3] Guido Van Rossum and Fred L. Drake. Python 3 Reference Manual. CreateSpace, Scotts Valley, CA, 2009.